

# RUBCRNCTCS

## Core

### Part 2



# Replication Control

A single pod is running an application- if it fails, it also crashes.

To prevent this, we want more than one POD running. Replication Control helps run multiple instances of a single POD in the Kubernetes cluster- providing high availability.

- Can be run with only one POD.
- Make sure there is a specific number of PODs up.

Also, it can create multiple PODs and share the load across them.

# Create Replication Controller

\$ **kubectl create -f rc-definition.yaml**

```
● ● ●  
1  apiVersion: v1  
2  kind: ReplicationController  
3  metadata:  
4    name: myapp-rc  
5    labels:  
6      app: my-app  
7      type: front-end  
8  
9  spec:  
10   template:  
11  
12     metadata:  
13       name: myapp-rc  
14       labels:  
15         app: my-app  
16         type: front-end  
17     spec:  
18       containers:  
19         - name: nginx-container  
20           image: nginx:1.15.5  
21   replicas: 3  
22
```

## View created RC

\$ **kubectl get replication controller**

# Replica Set

```
● ● ●

1  apiVersion: apps/v1
2  kind: Replication
3  metadata:
4    name: myapp-replicaset
5    labels:
6      app: myapp
7      type: front-end
8
9  spec:
10   template:
11     metadata:
12       name: myapp-pod
13       labels:
14         app: myapp
15         type: front-end
16     spec:
17       containers:
18         - name: nginx-container
19           image: nginx
20
21   replicas: 3
22
23   selector:
24     matchLabels:
25       app: myapp
26       type: front-end
27
```

Create Replication Set

\$ **kubectl create -f replicaset-definition.yaml**

# Deployment

Update, scale and rollback the enviornment



```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: my-app-deployment
5   labels:
6     app: myapp
7     type: front-end
8 spec:
9   template:
10    metadata:
11      name: myapp-pod
12      labels:
13        app: myapp
14        type: front-end
15    spec:
16      containers:
17        - name: nginx-container
18          image: nginx
19
20 replicas: 3
21 selector:
22   matchLabels:
23     type: front-end
```

# Services

Enable communication between various components within and outside of the app.

Help to connect other apps or users

Connect the Front-end to the users

Connect the Back-end to the Front-end pods

Connect to external data source

## 3 main services

NodePort

ClusterIP – Default

LoadBalancer

See all services

**\$ kubectl get services**

# NodePort

Listen to a port on the node and forward the request on that port, to a port in a POD running the app.

Multiple PODs in service- use the same selector

Multiple PODs in PODs- use the same labels

Also, work on multiple PODs- different nodes with the same ports- Auto

```
● ● ●  
1  apiVersion: v1  
2  kind: Service  
3  metadata:  
4    name: my-app-service  
5  
6  spec:  
7    type: NodePort  
8    ports:  
9      - targetPort: 8080  
10     port: 80  
11     nodePort: 30000 # Range 30000-32767  
12    selector:  
13      app: my-app  
14      type: front-end
```

# ClusterIP

The front end communicates with the back end and itself with Redis services.

All nodes have a non-static IP.

Kubernetes helps us to provide a single interface to access the pods in a group.

- A Service created for Back-end PODs will help group all the Back-end PODs together and provide a single interface for other PODs to access this service- the request is sent randomly.
- Each service get an IP and a name.

# ClusterIP “.yaml” file

```
● ● ●  
1  apiVersion: v1  
2  kind: Service  
3  metadata:  
4    name: my-app-service  
5  
6  spec:  
7    type: ClusterIP  
8    ports:  
9      - targetPort: 80  
10        port: 80  
11  
12    selector:  
13      app: my-app  
14      type: front-end  
15
```

# LoadBalancer

Can use a cloud load balancer, such as AWS, Google, and Azure.

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: my-app-service
5
6  spec:
7    type: LoadBalancer
8    ports:
9      - targetPort: 80
10        port: 80
11        nodePort: 30008
12
```

# Namespaces

- PODs, clusters, services, and nodes are created in a namespace by default.
- Kubernetes creates a set of PODs and services for internal purposes, such as those required by networking solutions, the DNS service, etc.
- To isolate the user and prevent you from deleting/modifying the services and create them in other cluster startup name kube-system.
- Resources that should be made available to all users are created in kube-public.
- Policies can be assigned to namespaces for permissions purposes.
- Can assign ‘quota’ of resources for each namespace- each namespace grants a certain amount of nodes and won’t use more than it’s allowed.

- Resources in the same namespace refer to each other with names.
- In other namespaces they refer as:  
**<service-name>.<namespace>.<service>.<domain>**
- Can view PODs in other namespace:  
**\$ kubectl get pods --namespace=<namespace-name>**
- Can create PODs in other namespace:  
**\$ kubectl create -f file.yaml --namespace=<namespace-name>**

# Create Namespace

- <file>.yaml



- Command:

**\$ kubectl create namespace <name>**

- Switch default namespace:

**\$ kubectl config set-context \$(kubectl current-context) --namespace=<name>**

# Imperative vs Declarative

Imperative: Step by step.

- A list of commands.
- To create use: ‘run’, ‘create’, and ‘expose’.
- To update: ‘edit’, ‘scale’, ‘replace’, and ‘set’.

Declarative: End goal

- Uses the ‘apply’ command.

**\$ kubectl apply**

- Take into consideration: the local configuration file, the live object definition on Kubernetes, and the last applied configuration.
- When using the command the <file>.yaml transform to <file>.json and store the last applied configuration.